

Name: \_\_\_\_\_  
Student ID: \_\_\_\_\_

## CS230 System Programming

### Midterm Exam

Tuesday, October 25, 2011 1:00-3:00pm

#### Instructions:

- Make sure that your exam is not missing any sheets, then write your full name on the front.
- Write your answers in the space provided below the problem. If you make a mess, clearly indicate your final answer.
- The exam has a maximum score of 90 points.
- This exam is CLOSED BOOK, CLOSED NOTES. You may only have two sheets of paper or notes with you. Good luck!

1 (12):
2 (13):
3 (12):
4 (10):
5 (8):
6 (15):
7 (10):
8 (10):
TOTAL (90):

**Problem 1. (12 points):**

Consider the following 8-bit floating point representation based on the IEEE floating point format:

- There is a sign bit in the most significant bit.
- The next 3 bits are the exponent. The exponent bias is  $2^{3-1} - 1 = 3$ .
- The last 4 bits are the fraction.
- The representation encodes numbers of the form:  $V = (-1)^s \times M \times 2^E$ , where  $M$  is the significand and  $E$  is the biased exponent.

The rules are like those in the IEEE standard(normalized, denormalized, representation of 0, infinity, and NAN). FILL in the table below. Here are the instructions for each field:

- **Binary:** The 8 bit binary representation.
- **M:** The value of the significand. This should be a number of the form  $x$  or  $\frac{x}{y}$ , where  $x$  is an integer, and  $y$  is an integral power of 2. Examples include 0,  $\frac{3}{4}$ .
- **E:** The integer value of the exponent.
- **Value:** The numeric value represented.

Note: you need not fill in entries marked with "—".

Description	Binary	$M$	$E$	Value
Minus zero				-0.0
—	0 100 0101			
Smallest denormalized (negative)				
Largest normalized (positive)				
One				1.0
—				5.5
Positive infinity		—	—	$+\infty$

**Problem 2. (13 points):**

Consider the source code below, where M and N are constants declared with `#define`.

```
int mat1[M][N];
int mat2[N][M];

int sum_element(int i, int j)
{
    return mat1[i][j] + mat2[i][j];
}
```

A. Suppose the above code generates the following assembly code:

```
sum_element:
    pushl %ebp
    movl %esp,%ebp
    movl 8(%ebp),%eax
    movl 12(%ebp),%ecx
    sall $2,%ecx
    leal 0(,%eax,8),%edx
    subl %eax,%edx
    leal (%eax,%eax,4),%eax
    movl mat2(%ecx,%eax,4),%eax
    addl mat1(%ecx,%edx,4),%eax
    movl %ebp,%esp
    popl %ebp
    ret
```

What are the values of M and N?

M =

N =

### Problem 3. (12 points):

Consider the following assembly code for a C for loop:

```
loop:
    pushl %ebp
    movl %esp,%ebp
    movl 8(%ebp),%ecx
    movl 12(%ebp),%edx
    xorl %eax,%eax
    cmpl %edx,%ecx
    jle .L4
.L6:
    decl %ecx
    incl %edx
    incl %eax
    cmpl %edx,%ecx
    jg .L6
.L4:
    incl %eax
    movl %ebp,%esp
    popl %ebp
    ret
```

Based on the assembly code above, fill in the blanks below in its corresponding C source code. (Note: you may only use the symbolic variables `x`, `y`, and `result` in your expressions below — *do not use register names*.)

```
int loop(int x, int y)
{
    int result;

    for ( _____; _____; result++ ) {
        _____;
        _____;
    }

    _____;

    return result;
}
```

#### Problem 4. (10 points):

Match each of the assembler routines on the left with the equivalent C function on the right.

foo1:

```
    pushl %ebp
    movl %esp,%ebp
    movl 8(%ebp),%eax
    sall $4,%eax
    subl 8(%ebp),%eax
    movl %ebp,%esp
    popl %ebp
    ret
```

foo2:

```
    pushl %ebp
    movl %esp,%ebp
    movl 8(%ebp),%eax
    testl %eax,%eax
    jge .L4
    addl $15,%eax
.L4:
    sarl $4,%eax
    movl %ebp,%esp
    popl %ebp
    ret
```

foo3:

```
    pushl %ebp
    movl %esp,%ebp
    movl 8(%ebp),%eax
    shrl $31,%eax
    movl %ebp,%esp
    popl %ebp
    ret
```

```
int choice1(int x)
{
    return (x < 0);
}
```

```
int choice2(int x)
{
    return (x << 31) & 1;
}
```

```
int choice3(int x)
{
    return 15 * x;
}
```

```
int choice4(int x)
{
    return (x + 15) / 4;
}
```

```
int choice5(int x)
{
    return x / 16;
}
```

```
int choice6(int x)
{
    return (x >> 31);
}
```

**Fill in your answers here:**

foo1 corresponds to choice \_\_\_\_\_.

foo2 corresponds to choice \_\_\_\_\_.

foo3 corresponds to choice \_\_\_\_\_.

The next problem concerns the following C code:

```
/* copy string x to buf */
void foo(char *x) {
    int buf[1];
    strcpy((char *)buf, x);
}

void callfoo() {
    foo("abcdefghi");
}
```

Here is the corresponding machine code on a Linux/x86 machine:

```
080484f4 <foo>:
080484f4: 55                pushl   %ebp
080484f5: 89 e5            movl    %esp,%ebp
080484f7: 83 ec 18        subl    $0x18,%esp
080484fa: 8b 45 08        movl    0x8(%ebp),%eax
080484fd: 83 c4 f8        addl    $0xffffffff8,%esp
08048500: 50              pushl   %eax
08048501: 8d 45 fc        leal    0xffffffffc(%ebp),%eax
08048504: 50              pushl   %eax
08048505: e8 ba fe ff ff  call    80483c4 <strcpy>
0804850a: 89 ec            movl    %ebp,%esp
0804850c: 5d              popl    %ebp
0804850d: c3              ret

08048510 <callfoo>:
08048510: 55                pushl   %ebp
08048511: 89 e5            movl    %esp,%ebp
08048513: 83 ec 08        subl    $0x8,%esp
08048516: 83 c4 f4        addl    $0xffffffff4,%esp
08048519: 68 9c 85 04 08  pushl   $0x804859c # push string address
0804851e: e8 d1 ff ff ff  call    80484f4 <foo>
08048523: 89 ec            movl    %ebp,%esp
08048525: 5d              popl    %ebp
08048526: c3              ret
```

### Problem 5. (8 points):

This problem tests your understanding of the stack discipline and byte ordering. Here are some notes to help you work the problem:

- `strcpy(char *dst, char *src)` copies the string at address `src` (including the terminating `'\0'` character) to address `dst`. It does **not** check the size of the destination buffer.
- Recall that Linux/x86 machines are Little Endian.
- You will need to know the hex values of the following characters:

Character	Hex value	Character	Hex value
'a'	0x61	'f'	0x66
'b'	0x62	'g'	0x67
'c'	0x63	'h'	0x68
'd'	0x64	'i'	0x69
'e'	0x65	'\0'	0x00

Now consider what happens on a Linux/x86 machine when `callfoo` calls `foo` with the input string “abcdefghi”.

- A. List the contents of the following memory locations immediately after `strcpy` returns to `foo`. Each answer should be an unsigned 4-byte integer expressed as 8 hex digits.

`buf[0]` = 0x\_\_\_\_\_

`buf[1]` = 0x\_\_\_\_\_

`buf[2]` = 0x\_\_\_\_\_

- B. Immediately **before** the `ret` instruction at address `0x0804850d` executes, what is the value of the frame pointer register `%ebp`?

`%ebp` = 0x\_\_\_\_\_

- C. Immediately **after** the `ret` instruction at address `0x0804850d` executes, what is the value of the program counter register `%eip`?

`%eip` = 0x\_\_\_\_\_

**Problem 6. (15 points):**

Consider the following incomplete definition of a C struct along with the incomplete code for a function func given below.

```
typedef struct node {
    _____ x;
    _____ y;
    struct node *next;
    struct node *prev;
} node_t;

node_t n;

void func() {
    node_t *m;
    m = _____;
    m->y /= 16;
    return;
}
```

When this C code was compiled on an IA-32 machine running Linux, the following assembly code was generated for function func.

```
func:
    pushl %ebp
    movl n+12,%eax
    movl 16(%eax),%eax
    movl %esp,%ebp
    movl %ebp,%esp
    shrw $4,8(%eax)
    popl %ebp
    ret
```

Given these code fragments, fill in the blanks in the C code given above. Note that there is a unique answer.

The types must be chosen from the following table, assuming the sizes and alignment given.

Type	Size (bytes)	Alignment (bytes)
char	1	1
short	2	2
unsigned short	2	2
int	4	4
unsigned int	4	4
double	8	4



## Problem 7. (10 points):

The following problem concerns basic cache lookups.

- The memory is byte addressable.
- Memory accesses are to **1-byte words** (not 4-byte words).
- Physical addresses are 13 bits wide.
- The cache is 2-way set associative, with a 4 byte line size and 16 total lines.

In the following tables, **all numbers are given in hexadecimal**. The contents of the cache are as follows:

2-way Set Associative Cache												
Index	Tag	Valid	Byte 0	Byte 1	Byte 2	Byte 3	Tag	Valid	Byte 0	Byte 1	Byte 2	Byte 3
0	09	1	86	30	3F	10	00	0	99	04	03	48
1	45	1	60	4F	E0	23	38	1	00	BC	0B	37
2	EB	0	2F	81	FD	09	0B	0	8F	E2	05	BD
3	06	0	3D	94	9B	F7	32	1	12	08	7B	AD
4	C7	1	06	78	07	C5	05	1	40	67	C2	3B
5	71	1	0B	DE	18	4B	6E	0	B0	39	D3	F7
6	91	1	A0	B7	26	2D	F0	0	0C	71	40	10
7	46	0	B1	0A	32	0F	DE	1	12	C0	88	37

## Part 1

The box below shows the format of a physical address. Indicate (by labeling the diagram) the fields that would be used to determine the following:

*CO* The block offset within the cache line

*CI* The cache index

*CT* The cache tag

12	11	10	9	8	7	6	5	4	3	2	1	0

## Part 2

For the given physical address, indicate the cache entry accessed and the cache byte value returned **in hex**. Indicate whether a cache miss occurs.

If there is a cache miss, enter “-” for “Cache Byte returned”.

**Physical address:** 0E34

A. Physical address format (one bit per box)

12	11	10	9	8	7	6	5	4	3	2	1	0
<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>

B. Physical memory reference

Parameter	Value
Byte offset	0x
Cache Index	0x
Cache Tag	0x
Cache Hit? (Y/N)	
Cache Byte returned	0x

**Problem 8. (10 points):**

After watching the mayor election, you decide to start a business in developing software for electronic voting. The software will run on a machine with a 1024-byte direct-mapped data cache with 64 byte blocks (i.e., cache line size is 64 bytes).

You are implementing a prototype of your software that assumes that there are 7 candidates. The C-structures you are using are:

```
struct vote {
    int candidates[7];
    int valid;
};

struct vote vote_array[16][16];
register int i, j, k;
```

You have to decide between two alternative implementations of the routine that initializes the array `vote_array`. You want to choose the one with the better cache performance.

You can assume:

- `sizeof(int) = 4`
- `vote_array` begins at memory address 0
- The cache is initially empty.
- The only memory accesses are to the entries of the array `vote_array`. Variables `i`, `j` and `k` are stored in registers.

A. What percentage of the writes in the following code will miss in the cache?

```
for (i=0; i<16; i++){
    for (j=0; j<16; j++) {
        vote_array[i][j].valid=0;
    }
}

for (i=0; i<16; i++){
    for (j=0; j<16; j++) {
        for (k=0; k<7; k++) {
            vote_array[i][j].candidates[k] = 0;
        }
    }
}
```

Total number of misses in the first loop: \_\_\_\_\_ %

Total number of misses in the second loop: \_\_\_\_\_ %

Overall miss rate for writes to vote\_array: \_\_\_\_\_ %

B. What percentage of the writes in the following code will miss in the cache?

```
for (i=0; i<16; i++){
    for (j=0; j<16; j++) {
        for (k=0; k<7; k++) {
            vote_array[i][j].candidates[k] = 0;
        }
        vote_array[i][j].valid=0;
    }
}
```

Miss rate for writes to vote\_array: \_\_\_\_\_ %